

# **IFT 6135: Representation Learning**

Michael Noukhovitch

Winter 2019

Notes written from Aaron Courville's lectures.

# Contents

<b>1</b>	<b>Neural Networks</b>	<b>3</b>
1.1	Artificial Neuron . . . . .	3
1.2	Activation Functions . . . . .	3
1.3	Neural Networks . . . . .	3
1.3.1	Single Layer . . . . .	3
1.3.2	Multi Layer . . . . .	3
1.4	Biological Inspiration . . . . .	4
<b>2</b>	<b>Training Neural Networks</b>	<b>4</b>
2.1	Empirical Risk Minimization . . . . .	4
2.2	Stochastic Gradient Descent . . . . .	4
2.3	Gradient Computation . . . . .	4
2.4	Backpropagation . . . . .	5
2.5	Flow Graph . . . . .	5

# 1 Neural Networks

## 1.1 Artificial Neuron

$g(b + w^T x)$

**pre-activation**  $b + w^T x$

**connection weights**  $w$

**neuron bias**  $b$

**activation function**  $g$

## 1.2 Activation Functions

**linear**  $a$

**sigmoid**  $\frac{1}{1 + \exp(-a)}$

**tanh**  $\frac{\exp a - \exp -a}{\exp(a) + \exp(-a)}$

**ReLU**  $\max(0, a)$

**maxout**  $\max_{j \in [1, k]} a_j$

**softmax**  $\frac{\exp(a_i)}{\sum_c \exp(a_c)} \forall i$

## 1.3 Neural Networks

### 1.3.1 Single Layer

**neuron capacity** a single neuron can do binary classification iff linearly separable

**universal approximation theorem** (Hornik, 1991) a single-layer NN can approximate any continuous function given enough hidden units

### 1.3.2 Multi Layer

**input**  $h^{(0)} = x$

**hidden layer pre**  $a^{(k)}(x) = b^{(k)} + W^{(k)} h^{(k-1)}(x)$

**hidden layer activation**  $h^{(k)} = g(a^{(k)}(x))$

**output**  $h^{(L+1)}(x) = o(a^{(L+1)}(x))$

## 1.4 Biological Inspiration

# 2 Training Neural Networks

## 2.1 Empirical Risk Minimization

learning as optimization

$$\operatorname{argmin}_{\theta} \frac{1}{T} \sum_t l(f(x^{(t)}; \theta), y^{(t)}) + \lambda \Omega(\theta)$$

**loss function**  $l$  is a surrogate for what we truly want (upper bound)

**regularizer**  $\Omega$  penalizes certain values of  $\theta$

## 2.2 Stochastic Gradient Descent

```
initialize  $\theta$ 
for  $n$  epochs do
    foreach training example  $x^{(t)}, y^{(t)}$  do
         $\Delta \leftarrow -\nabla_{\theta} l(f(x^{(t)}; \theta), y^{(t)}) - \lambda \nabla_{\theta} \Omega(\theta)$ 
         $\theta \leftarrow \theta + \alpha \Delta$ 
    end
end
```

where  $\nabla$  is the gradient,  $\alpha$  is the learning rate

## 2.3 Gradient Computation

given a categorical output with classes  $c$ , let softmax output be  $f(x)_c = p(y = c|x)$

gradients for output

$$\begin{aligned} \frac{\partial}{\partial f(x)_c} - \log f(x)_y &= \frac{-1_{(y=c)}}{f(x)_y} \\ \nabla_{f(x)} - \log f(x)_y &= \frac{-1}{f(x)_y} \begin{bmatrix} 1_{(y=0)} \\ \dots \\ 1_{(y=C-1)} \end{bmatrix} \\ &= \frac{-e(y)}{f(x)_y} \end{aligned}$$

gradients for pre-activation

$$\begin{aligned} \frac{\partial}{\partial a^{(L+1)}(x)} - \log \sigma(a^{(L+1)}(x))_y &= -(1_{(y=c)} - f(x)_y) \\ \nabla_{f(x)} - \log f(x)_y &= -(e(y) - f(x)) \end{aligned}$$

where  $e(y)$  gives the one-hot vector of length  $C$  with 1 at index  $y$ , and  $\sigma$  is the sigmoid function

## 2.4 Backpropagation

To simplify things, use the chain rule to rewrite gradients in terms of in terms of the layers above them

```

compute output gradient  $\nabla_{a^{(L+1)}(x)} - \log f(x)_y = -(e(y) - f(x))$ 
for  $k = L + 1 \rightarrow 1$  do
  hidden layer weights
     $\nabla_{W^{(k)}(x)} - \log f(x)_y = (\nabla_{a^{(k)}(x)} - \log f(x)_y) h^{(k-1)}(x)^T$ 
  hidden layer biases
     $\nabla_{b^{(k)}(x)} - \log f(x)_y = \nabla_{a^{(k)}(x)} - \log f(x)_y$ 
  output below
     $\nabla_{h^{(k-1)}(x)} - \log f(x)_y = W^{(k)T} (\nabla_{a^{(k)}(x)} - \log f(x)_y)$ 
  pre-activation below
     $\nabla_{a^{(k-1)}(x)} - \log f(x)_y = (\nabla_{h^{(k-1)}(x)} - \log f(x)_y) \odot [\dots, g'(a^{(k-1)}(x)_j), \dots]$ 
end

```

## 2.5 Flow Graph