

SE 350: Operating Systems

Michael Noukhovitch

Winter 2015, University of Waterloo

Notes written from Thomas Reidemeister's lectures.

Contents

1	Introduction	4
1.1	Definitions	4
1.2	Interrupts	4
1.2.1	Type of interrupts	4
1.2.2	How interrupts work	4
1.3	Multiple interrupts	5
1.3.1	Sequential	5
1.3.2	Nested	5
1.4	Peripheral interrupts	5
1.4.1	Programmed I/O	5
1.4.2	Interrupt-driven I/O	5
1.4.3	Direct memory access	5
1.5	Memory hierarchy	5
1.5.1	Hierarchy	5
1.5.2	Cache	6
2	Operating Systems Overview	6
2.1	Definition	6
2.2	OS Innovations	6
2.2.1	Hardware Features	6
2.2.2	Modes of operation	6
2.2.3	Multiprogramming	7
2.2.4	Time Sharing	7
2.3	Major Achievements	7
2.3.1	Processes	7
2.3.2	Memory Management	7
2.3.3	Information Security	7
2.3.4	Scheduling and Resource Management	8
2.3.5	System Structure	8
2.4	Modern Operating System	8
3	Processes	9
3.1	Process Elements	9
3.2	Process Model	9
3.2.1	Five State Model	9
3.2.2	Process Creation	9
3.2.3	Process Termination	9
3.2.4	Process Suspension	10
3.2.5	Process Switching	10
3.3	OS Control Structures	10
3.3.1	Memory Tables	10
3.3.2	I/O Tables	10
3.3.3	File Tables	11
3.3.4	Process Table	11

4	Threads, SMP, Microkernels	11
4.1	Threads	11
4.1.1	Benefits	11
4.1.2	States	11
4.1.3	Approaches	11
4.2	Microkernel	12
4.2.1	Benefits	12
4.2.2	Design	12
5	Concurrency	12
5.1	Terms	12
5.2	Hardware Support	13
5.2.1	Interrupt Disabling	13
5.2.2	Machine Instructions	13
5.3	Semaphore	13
5.3.1	Operations	14
5.3.2	Producer-Consumer	14
5.4	Monitor	15
5.5	Message Passing	15
5.5.1	Synchronization	15
5.5.2	Addressing	15

1 Introduction

1.1 Definitions

Operating Systems: a standardized abstraction from hardware that:

- manages resources
- provides set of services
- consumes resources

Instruction execution:

1. fetches instruction into IR
2. executes instruction

In reality, it is a bit more complicated: there is a pipeline, out of order execution...

1.2 Interrupts

1.2.1 Type of interrupts

- **program** result of instruction execution e.g. arithmetic overflow
- **timer** timer within process, allows OS to perform regular functions
- **I/O** generated by I/O controller
- **hardware failure** power failure or memory parity failure

1.2.2 How interrupts work

Hardware:

1. interrupt issued
2. processor finishes current instruction
3. acknowledge interrupt
4. push PSW and PC onto control stack
5. load new PC

Software:

6. save remainder of process state
7. interrupt
8. restore process state information
9. restore PSW and PC

1.3 Multiple interrupts

Two ways of handling an interrupt during an interrupt:

1.3.1 Sequential

Ignore any interrupts when you are in an interrupt and when done, check for interrupts that occurred.

1.3.2 Nested

If the second interrupt is of higher priority, recurse into it. Otherwise, wait until interrupt is finished.

1.4 Peripheral interrupts

Peripherals such as hard drives take a while to complete their action. As opposed to waiting and wasting that time, we use an interrupt to execute other instructions while our I/O process runs.

1.4.1 Programmed I/O

No interrupts occur, just wait until the I/O is complete

1.4.2 Interrupt-driven I/O

Processor interrupted when I/O is ready and all writes/reads are passed through CPU into memory. This is faster, since there is no waiting.

1.4.3 Direct memory access

Transfers a block of data directly into memory and interrupt is sent when process is complete. This is more efficient because data does not need to go through CPU. Not always available (e.g. external peripheral)

1.5 Memory hierarchy

Major constraints in memory:

- size
- speed
- cost

1.5.1 Hierarchy

Top: Inboard memory

Middle: Outboard storage

Bottom: Offline storage

1.5.2 Cache

Cache: small, fast memory, invisible to the OS, that speeds up accesses exploiting the principle of locality

2 Operating Systems Overview

2.1 Definition

Operating System a program that controls the execution of applications and is a standard interface between hardware and software

- **convenience:** need no knowledge of hardware
- **efficiency:** move optimization from devs to tools
- **ability to evolve:** can replace internals

Kernel: portion of the OS in main memory; a nucleus that contains frequently used functions
OS services

- program development and execution
- access & control of access to I/O
- system access control
- error detection and response
- accounting

2.2 OS Innovations

2.2.1 Hardware Features

- **Memory protection:** do not allow memory containing monitor to be altered
- **Timer:** prevents a job from monopolizing system
- **Privileged instruction:** certain instructions can only be executed by the monitor (e.g. I/O)
- **Interrupts**

2.2.2 Modes of operation

To protect users from each other (and the kernel from user), we have two modes:

- **User mode:** not privileged
- **Kernel mode:** privileged and access to protected memory

2.2.3 Multiprogramming

When one job needs to wait for I/O, the processor can switch to another job. The timer is also used to switch processes and stop monopolization.

- maximize processor use
- use job control language

2.2.4 Time Sharing

Processor time is shared by multiple users

- minimize response time

2.3 Major Achievements

2.3.1 Processes

Process: a program in execution

- a program
- associated data
- execution content (needed by OS)

2.3.2 Memory Management

- process isolation
- automatic allocation
 - virtual memory: allows programmer to address memory without regard to physical addressing
 - paging: allows processes to be comprised of fixed-size blocks (pages)
- swap program code
- shared memory
- long term storage

2.3.3 Information Security

- availability: protecting system against interruption (downtime)
- confidentiality: authorizing data (chmod)
- data integrity: protect from modification
- authenticity: verifying identity of users

2.3.4 Scheduling and Resource Management

- fairness: give equal access to resources
- differential responsiveness: discriminate by class of jobs
- efficiency: maximize throughput

2.3.5 System Structure

The system as a hierarchical structure with each level relying on lower levels for its functions.

1. circuits: registers, gates, buffers
2. instructions: add, subtract, load, store
3. procedures: call stack, subroutine
4. interrupts
5. processes: suspend, wait, resume
6. local store: blocks of data, allocate
7. virtual memory: segments, pages
8. communications: pipes
9. file system: files
10. external devices
11. directories
12. user process
13. shell

2.4 Modern Operating System

Developments leading to modern operating systems:

- **Microkernel architecture:** only essentials to kernel, everything else in user space (e.g. QNX)
- **Multithreading:** process divided into concurrent threads
- **Symmetric multiprocessing:** multiple processors share main memory and I/O
- **Distributed OS:** illusion of a single main and secondary memory
- **Asymmetric multiprocessing:** one big processor controls many small ones
- **Object oriented design:** customize OS without disrupting system

3 Processes

3.1 Process Elements

Process Control Block (PCB): data structure that contains process elements

- **identifier**
- **state**
- **priority**
- **memory pointers**
- **context data:** registers, PSW, PC
- **I/O status information**
- **accounting information:** processor time, time limits, threads

3.2 Process Model

3.2.1 Five State Model

- **running:** currently executing
- **ready:** can be executed
- **waiting:** can't execute, blocked by something
- **new**
- **exit:** halted or aborted

Since the processor is faster than I/O, we may sometimes want to admit more processes even after we are out of memory. To do this we **swap** out processes to disk that are waiting and we get two new states:

- **blocked/suspend**
- **ready/suspend**

3.2.2 Process Creation

- new batch job
- interactive logon
- created by OS for a service
- created by existing process

3.2.3 Process Termination

- normal completion
- time limit exceeded
- time overrun
- memory unavailable
- bounds error (segfault)
- protection error
- arithmetic error
- I/O error
- invalid instruction
- privileged instruction
- data misuse
- OS intervention
- parent termination
- parent request

3.2.4 Process Suspension

- swapping
- other OS reason
- user request
- timing
- parent request

3.2.5 Process Switching

clock interrupt : maximum allowed time surpasses

I/O interrupt : I/O completed

memory fault : memory address is in virtual memory, bring into main memory

trap : error or exception, used for debugging

supervisor call : switch to kernel process

3.3 OS Control Structures

3.3.1 Memory Tables

- keeps track of main and secondary memory
- protection for access to shared memory
- information to manage virtual memory

3.3.2 I/O Tables

- manages I/O devices (status and availability)
- location in main memory for transferring I/O

3.3.3 File Tables

- manages existence and location of files
- attributes and status of files (e.g. rwxr...)

3.3.4 Process Table

- where process is located in memory as a **process image**:
 - program
 - data
 - system stack
 - PCB

4 Threads, SMP, Microkernels

4.1 Threads

Thread: execution entity under a process **Multithreading**: multiple threads of execution within a single process

4.1.1 Benefits

- less time to create/terminate (no kernel resource allocation)
- less time to switch between threads v. processes
- sharing memory between threads

4.1.2 States

- spawn
- block
- unblock
- finish

4.1.3 Approaches

User-level threads:

- thread management by application
- less switching overhead
- scheduling is app-specific
- can run on any OS

Kernel-level threads:

- thread management by kernel
- can schedule threads on multiple processors
- OS calls blocking only the thread

Combined approach:

- threads can be grouped to kernel threads
- kernel knows/schedules threads and processes

4.2 Microkernel

Microkernel: small operating system that only contains essential core functions

4.2.1 Benefits

- uniform interface on request by process
- extensibility
- flexibility
- portability
- reliability
- distributed systems support
- object-oriented OS

4.2.2 Design

- low-level memory management: map each virtual page to a physical page
- interprocess communication: copying messages
- I/O and interrupt management: interrupts as messages, no knowledge of IRQ handlers

5 Concurrency

5.1 Terms

critical section : section of code that may have concurrency issues (shared memory ...)

deadlock : processes are locked because each is waiting for the other

livelock : two processes cycle uselessly because of the others' cycling

mutual exclusion : shared resources may not be modified concurrently

race condition : multiple threads/processes read and write shared memory concurrently

starvation : a runnable process is overlooked indefinitely by the scheduler

For mutual exclusion requires:

- 1 process at a time (per resource) in the critical section
- process halting in non-critical section must not interfere with other processes
- no deadlock or starvation
- no delay to critical section if it is unblocked
- no assumptions about relative process speeds or quantities
- process remains inside critical section for finite number of time

and can be achieved by any of the following options:

5.2 Hardware Support

5.2.1 Interrupt Disabling

- guarantees mutual exclusion on uniprocessor
- no guarantee for mutliprocessor
- a process runs until interrupt or it invokes OS service

5.2.2 Machine Instructions

- performed in single instruction cycle
- access to memory is blocked for other instructions

Advantages:

- applicable to any number of processes (single processor or multiprocessor)
- simple and easy to verify
- can support multiple critical sections

Disadvantages:

- busy-waiting consumes time
- starvation possible (one proc leaves critical section and multiple waiting)
- deadlock possible (low priority holds critical section but OS switched to high priority)
- pipeline stalls

5.3 Semaphore

Semaphore: variable with integer value used for signaling

5.3.1 Operations

- initialize with non-negative value
- **wait** decrements value
- **signal** increments value
- only the process with the lock can release the lock

Example 5.1. Semaphore primitive

```
struct semaphore {
    int count;
    queueType queue;
}

void semWait(semaphore s) {
    s.count--;
    if (s.count < 0) {
        s.queue.push(thisProcess);
        block(thisProcess);
    }
}

void semSignal(semaphore s) {
    s.count++;
    if (s.count <= 0) {
        nextProcess = s.queue.pop();
        addToReady(nextProcess);
    }
}
```

5.3.2 Producer-Consumer

Concurrency problem where a producer adds elements to a buffer and a consumer takes them out, we can use extra semaphores to keep track of other critical values

Example 5.2. Infinite-buffer producer/consumer

```
semaphore n = 0;
semaphore s = 1;
void producer() {
    while (true) {
        produce();
        semWait(s);
        append();
        semSignal(s);
        semSignal(n);
    }
}

void consumer() {
    while (true) {
        semWait(n);
```

```

        semWait(s);
        take();
        semSignal(s);
        consume();
    }
}

```

5.3.3 Reader-Writer

Concurrency problem where many readers can read from a file, but only one writer may write to it and all readers must wait for the writer

Example 5.3. Readers/Writers problem with Reader priority

```

int readCount;
semaphore x = 0, wsem = 1;
void reader() {
    while (true) {
        semWait(x);
        readCount++;
        if (readCount == 1)
            semWait(wsem);
        semSignal(x);
        READUNIT();
        semWait(x);
        readCount--;
        if (readCount == 0)
            semSignal(wsem);
        semSignal(x);
    }
}
void writer() {
    while (true) {
        semWait(wsem);
        WRITEUNIT();
        semSignal(wsem);
    }
}

```

5.4 Monitor

Monitor: a software module that controls entry to data using a mutex

- uses condition variables for signaling
- unused signals are lost (*diff from semaphore*)

5.5 Message Passing

Enforce mutual exclusion by information exchange

5.5.1 Synchronization

Blocking : sender and receiver are blocked until message delivered (rendevous)

Non-blocking Send : only receiver blocked until message arrives

Non-blocking : no waiting

5.5.2 Addressing

Direct Addressing:

- send has a specific identifier of destination process
- receive knows which process to expect
- receive uses source parameter to return value after message received

Indirect Addressing:

- message sent to shared data structure made of queues (**mailboxes**)
- processes send messages to a mailbox for others to pick up

Example 5.4. Indirect messaging

```
const int n = /* number of processes */
create_mailbox(mutex);
void P (int i) {
    message msg;
    while (true) {
        receive(mutex, msg);
        /* critical section */
        send(mutex, msg);
        /* remainder */
    }
}
```